

Incident Response Standard Operating Procedure

This SOP provides the detailed sequence for declaring, coordinating, communicating, and resolving service incidents.

Department: Sop | Version: 1.1 | Owner: Bobby Hall, Site Reliability Manager

Access Level: Internal - Operations and Engineering | Created: 2024-07-03 | Updated: 2024-11-30 | Source File: sop_incident_response_v1.pdf

1.0 Purpose

The incident response SOP ensures that service degradation or outage events are handled using a common and time-aware operating model. In a production-grade RAG dataset, the document needs to answer both classification questions and urgent procedural questions, such as what happens first, who owns coordination, and when updates should be issued.

The procedure is designed to reduce confusion under pressure. By standardizing declaration, role assignment, mitigation, communication, and review, it helps teams stabilize service faster while creating the evidence needed for post-incident learning and governance.

2.0 Scope

This SOP applies to incidents affecting availability, data integrity, security posture, customer experience, or business-critical internal operations. It is used when an issue requires coordinated response across more than one person or function, even if the root cause is not yet known.

The procedure covers severity classification, incident declaration, immediate response, communication cadence, recovery, and post-incident review. Security incidents may require additional specialist procedures, but the coordination model here still applies unless explicitly superseded.

- Applies to customer-visible outages, severe degradation, and significant internal operational failures.
- Requires a named incident commander and a durable incident record for coordinated events.
- Works alongside security, vendor, or product-specific response procedures where needed.

3.0 Responsibilities

The responsibilities below define who must act, approve, review, and maintain records so the process can be executed consistently across teams and audit periods.

3.1 Incident Commander

- Own incident direction, assign roles, and keep the response focused on service restoration.
- Decide when to escalate severity, request additional responders, or move to rollback or mitigation.
- Ensure updates remain regular, factual, and tied to concrete next actions.

3.2 Technical Responders

- Investigate symptoms, test mitigations, and communicate findings clearly to the commander.
- Avoid competing uncoordinated changes during the incident unless explicitly assigned.
- Record meaningful technical observations and actions in the incident log.

3.3 Operations and Communications Partners

- Maintain the incident record, stakeholder updates, and timeline accuracy.
- Coordinate customer-facing messaging when external communication is needed.
- Ensure follow-up actions are assigned before closure.

4.0 Incident Declaration Procedure

The first step is to determine whether the issue requires formal incident handling. Teams should favor declaring early when customer impact or uncertainty is high, because a low-friction declaration is usually cheaper than a delayed coordinated response after the blast radius has widened.

4.1 Step 4.1 Classify Severity

- Assess user impact, revenue risk, security concern, and time sensitivity.
- Assign the initial severity even if some technical facts remain unclear.
- Update the severity later when better evidence becomes available.

4.2 Step 4.2 Open the Incident Record

Create the incident record and response channel within ten minutes of declaration. Include the initial summary, suspected systems affected, severity, incident commander, and the next expected update time so responders and stakeholders know where authoritative information will live.

- Open the designated response channel and tracking record.
- Assign incident commander and note current responder list.
- Capture the initial problem statement and expected next update.

5.0 Immediate Response Procedure

The response should prioritize stabilization before deep root-cause analysis. Early actions should aim to reduce user harm, restore essential functionality, and limit blast radius while enough signal is collected to support sound technical decisions.

5.1 Step 5.1 Establish Working Hypotheses

- Review alerts, dashboards, recent changes, and known dependencies.
- State likely causes as hypotheses and separate them from confirmed facts.
- Assign owners to parallel investigation tracks only when coordination can be maintained.

5.2 Step 5.2 Apply Mitigation

Use the lowest-risk mitigation available that can materially reduce impact, such as rollback, traffic shift, feature disablement, or queue protection. Mitigation steps should be logged with timestamps because they become critical context for both restoration and later analysis.

- Choose mitigations that reduce customer harm before pursuing perfection.
- Record every configuration or deployment change made during the response.
- Reassess severity and required responders after mitigation is applied.

6.0 Communication, Recovery, and Review Procedure

Clear communication is part of incident handling, not a separate administrative burden. Internal stakeholders need regular, time-stamped updates with current impact, actions taken, and next steps. When external communication is required, the designated communications lead should coordinate messaging with the incident commander so customers receive accurate and consistent information.

6.1 Step 6.1 Maintain Update Cadence

- Post initial summary, current impact, and next steps as soon as the incident is declared.
- Update every fifteen to thirty minutes for active high-severity incidents.
- Use time stamps and avoid speculation presented as fact.

6.2 Step 6.2 Confirm Recovery

- Verify customer impact has ended and health indicators remain stable.
- Watch for recurrence during the stabilization window before declaring resolution.
- Confirm dependent systems and backlog processing have returned to normal.

6.3 Step 6.3 Complete the Post-Incident Review

After recovery, the commander or designated owner must prepare a review covering timeline, root cause, contributing factors, customer impact, and preventive actions. The goal is not blame allocation; it is to improve systems, processes, and readiness for the next event.

- Assign owners and due dates for follow-up actions.
- Link evidence, dashboards, and decision points used during the response.
- Share the review with the relevant leadership and affected teams.

7.0 Exceptions and Escalation

If the response deviates from this SOP because of security handling, vendor involvement, or unavailable systems, the deviation should be recorded with a reason and the alternative coordination method used. Hidden process changes create confusion during review and repeat events.

Escalate quickly when incident impact expands, responders become overloaded, or communication quality drops. Coordination failure is itself an incident risk and should be treated as such.

8.0 Records and Review

The incident record must include timeline entries, severity changes, responder assignments, communications, mitigations, and closure evidence. Durable records matter because post-incident learning depends on facts captured during the event rather than reconstructed memory.

This SOP is reviewed after major incidents and at least semiannually. Repeated response pain points should trigger updates to tools, role definitions, or mitigation playbooks in addition to procedural wording.